

# Kyumin Han

+1 (856) 379-5682 | [hankyumin05@gmail.com](mailto:hankyumin05@gmail.com) | [linkedin.com/in/kyu-han-pitt2028](https://www.linkedin.com/in/kyu-han-pitt2028) | [github.com/kyuhan28](https://github.com/kyuhan28)

## EDUCATION

University of Pittsburgh, B.S. in Computer Science, Minor in Business

Expected Graduation: May 2028

GPA: 3.32 / 4.0

Pittsburgh, PA

**Relevant Coursework:** Data Structures & Algorithms, Software Engineering, Assembly Language, Systems Software

**Member:** Computer Science Club, AWS Cloud Club, Taekwondo Club

## TECHNICAL SKILLS

**Programming Languages:** Java, Python, C, HTML, CSS, C++, SQL, NoSQL

**Frameworks & Tools:** Pandas, React, Node.js, Git, GitHub, Jupyter, AWS, REST APIs, OpenRouter

## WORK EXPERIENCES

University of Pittsburgh

Pittsburgh, PA

Undergraduate Teaching Assistant

July 2026 - Present

- Host two open class labs for students to learn difficult topics about **Python** programming, topics, and ideas
- Hosting weekly meetings with professors to strengthen the course and creating weekly agendas for students.
- Held **6** hours of weekly office hours, mentoring students in Python to boost lab completion rates by **30%**
- Demonstrating leadership by teaching many students about various topics in programming, data structures, and algorithms.

## PROJECTS

**Multi-Lingual LLM Medical Research** | Python, Qwen3-8B, Hugging Face, Jupyter Notebook

March 2024 – July 2024

- **Developed** a model-agnostic uncertainty estimation pipeline in **Python** (with Google Colab) using **Qwen3-8B** to classify ~5,000 medical triage scenarios as non-emergent, partially emergent, or emergent across English, Hindi, and Romanized Hindi inputs.
- **Engineered** a **2S** (semantic similarity) **confidence score** to quantify model certainty on critical emergency predictions across datasets. Wanted to clarify if the model was confident in its answers before proceeding.
- **Evaluated** cross-lingual consistency using **BERTScore** precision and recall metrics across identical prompts in multiple languages. Knowing how precise the model is allows me to factor out language barriers.
- **Found** systemic model overconfidence, identifying the LLM produced a **>75% confidence score** on nearly all 5,000 cases, regardless of input language.

**NVIDIA Verification Layer** - SteelHacks XIII | Python, Jupyter Notebook, Nemotron Ultra and Super

Sept 2026

- **Built** a safety gate that audits LLM-generated medical advice: an LLM extracts findings, then a deterministic **Python** policy layer converts findings into a show/withhold decision, so every withheld answer names the rule that blocked it.
- Evaluated on a **41-case** labeled set built by injecting known flaws into correct drafts, making ground truth free. The gate caught **26/31** unsafe drafts while shipping **10/10** safe ones. The model's own risk score alone caught only **17/31** at identical coverage.
- A same-prompt GPT-4o-mini judge caught **30/31** but shipped only **2/10** safe answers - so recall alone can't distinguish a usable gate from one that blocks the product.

**Mental Awareness-Focused Interactive Experience** - SteelHacks XI | Python, React, Node.js

Sept 2024

- Developed an interactive story-game using **Python** to spread awareness on mental health.
- Created decisive trees and branching whilst incorporating user experience to generate more branches based on their decision.